

AI x TRADING: Strategic Research Dossier

The State of the Art — November 2025

EXECUTIVE SUMMARY

The intersection of AI and trading has reached an inflection point. The market for AI-driven trading platforms grew from \$3.21 billion in 2024 to a projected \$20.33 billion by 2034. Numerai's crowdsourced AI hedge fund delivered 25.45% net returns in 2024 and reached a \$500 million valuation. AI agent tokens collectively surged from \$4.8 billion to \$50.5 billion in market cap between mid-2024 and early 2025. This dossier provides a comprehensive landscape analysis across TradFi, DeFi, and prediction markets to inform your next-generation trading solution.

1. THE FRONTIER PROJECTS (The "Who")

1.1 AI AS ALPHA: Autonomous Agents Executing Trades

Tier 1: Institutionally-Validated

Numerai — The "Last Hedge Fund"

- **Status:** Series C at \$500M valuation (November 2025), backed by J.P. Morgan Asset Management (\$500M capacity commitment), Paul Tudor Jones, Union Square Ventures, and top university endowments
- **Performance:** 25.45% net return in 2024 (one down month), 8% net return through October 2025
- **AUM:** Grew from \$60M to \$550M in three years
- **Mechanism:** Global data science tournament where thousands of ML models contribute stock-picking signals aggregated into a "Meta Model"
- **Edge:** Crowdsourced intelligence + crypto-economic alignment via NMR token staking

Renaissance Technologies (Medallion Fund)

- **Status:** \$66B equity portfolio, 3,400+ holdings (as of Q1 2025)
- **Historical Performance:** 66% gross / 39% net annual returns since 1988
- **Note:** Closed to outside investors; serves as the benchmark for what's possible
- **Technical Approach:** Statistical arbitrage, high-frequency pattern recognition, massive compute for signal discovery

Tier 2: DeFi-Native Autonomous Agents

Giza Protocol (ARMA Agent)

- **Status:** \$20M in Assets Under Agents, \$5.4M transaction volume in 4 weeks on Base
- **Performance:** 83% increase in DeFi lending yields during volatile weekends; 15% USDC yields
- **Technical Stack:** Semantic abstraction layer + EigenLayer AVS for decentralized execution + smart account infrastructure
- **Edge:** Non-custodial autonomous yield optimization across Morpho, Aave, Moonwell

ElizaOS (formerly ai16z)

- **Status:** First AI-led VC DAO, \$28M AUM at peak, rebranded January 2025
- **Framework:** Eliza multi-agent simulation platform (GitHub's #2 trending repo, 2,800+ forks)
- **Mechanism:** "Marketplace of Trust" where token holders provide investment recommendations; AI agent assigns trust scores
- **Notable:** Raised 420.69 SOL initially; reached \$2B+ market cap

Virtuals Protocol

- **Status:** \$4.6B peak market cap (January 2025); 21,000+ AI agents launched
- **Ecosystem:** AIXBT (social sentiment AI), Luna (AI influencer), VADER
- **Platform:** AI agent launchpad on Base and Solana with tokenization and co-ownership
- **Edge:** GAME framework (Generative Autonomous Multimodal Entities) for autonomous agent behavior

Tier 3: Retail/Emerging Platforms

Platform	Key Feature	Status
3Commas	SmartTrade Terminal, SmartSell AI	Industry leader by user base
Bitsgap	Grid trading + AI optimization	Paid tool, multi-exchange
Coinrule	Rule-based strategy creation	Strategy creators focus
Cryptohopper	Cloud-based, sentiment analysis	Cross-exchange automation
WunderTrading	Statistical arbitrage bots	Portfolio management

1.2 AI AS ANALYST: LLMs for Sentiment & Signal Generation

Research-Validated LLM Performance

FinLlama/FinDPO Framework

- **Model:** Fine-tuned Llama 2 7B for financial sentiment

- **Performance:** 67% annual returns with Sharpe ratio of 2.0 (after 5bps costs)
- **Advantage:** 11% improvement over traditional supervised fine-tuning

OPT (GPT-3 Based) for Financial News

- **Dataset:** 965,375 U.S. financial news articles (2010-2023)
- **Accuracy:** 74.4% prediction accuracy for stock returns
- **Strategy Performance:** Sharpe ratio of 3.05 (vs 1.23 for dictionary models)
- **Result:** 355% gains August 2021 - July 2023

AIXBT Social Intelligence Agent

- **Function:** Monitors 400+ crypto KOLs, provides real-time market narratives
- **Claimed Performance:** 48% win rate, 19% average returns across 416 token calls
- **Caveat:** Accuracy contested (some analyses show ~31% actionable signals)
- **Access:** 600,000 AIXBT tokens (~\$200K) required for terminal access

Open-Source Financial LLMs

Model	Base	Purpose	Notable
FinGPT	Various LLMs	Sentiment, robo-advising	<\$300 fine-tuning cost
FinBERT	BERT 110M	Financial sentiment	Financial lexicon optimized
BloombergGPT	50B params	Financial NLP	Proprietary data access
FinGPT-Forecaster	Llama variants	Stock price prediction	HuggingFace available

1.3 AI AS INFRASTRUCTURE: Decentralized Compute & Execution

Decentralized AI Networks

Bittensor (TAO)

- **Market Cap:** ~\$3B (November 2025)
- **Mechanism:** Proof-of-intelligence consensus; 120+ specialized subnets
- **Function:** Decentralized ML model marketplace; validators train/rank models
- **Recent:** Dynamic TAO upgrade (Feb 2025) shifted to performance-based rewards
- **Institutional Interest:** Nasdaq-listed companies acquired \$17.5M TAO since June 2025

Artificial Superintelligence Alliance (FET)

- **Components:** Fetch.ai, SingularityNET, Ocean Protocol, CUDOS merger

- **Platform:** Agentverse marketplace for autonomous agent deployment
- **Notable:** ASI-1 Mini Web3-native LLM; ASI Chain in development

Prediction Market Infrastructure

Polymarket

- **Status:** \$9B valuation (ICE invested \$2B in November 2025)
- **Volume:** \$1.5B+ monthly (2025)
- **Edge:** UMA Protocol optimistic oracle for settlement; CLOB hybrid matching
- **AI Integration:** Polymarket Agents repo on GitHub for autonomous trading

Kalshi

- **Status:** \$5B valuation; \$50B annualized volume (2025)
- **Regulatory:** CFTC-approved; >62% U.S. market share
- **Use Case:** Institutional-grade prediction market access

2. THE TECHNICAL "SECRET SAUCE" (The "How")

2.1 Reinforcement Learning (RL) Approaches

Proven Algorithms

Algorithm	Use Case	Performance Notes
PPO (Proximal Policy Optimization)	Portfolio management	Stable training, policy constraints
A2C (Advantage Actor-Critic)	Multi-asset trading	Good for discrete actions
TD3 (Twin Delayed DDPG)	Continuous action spaces	Dynamic share sizing
SAC (Soft Actor-Critic)	Risk-sensitive trading	Entropy regularization
DQN (Deep Q-Network)	Single-asset trading	Outperforms SMA/buy-hold

Key Research Finding: A study on reinforcement learning in crypto trading found AI agents using self-created strategies delivered profits of 9.94% to 31.53%, compared to 8.33% average for traditional bots.

Ensemble Methods

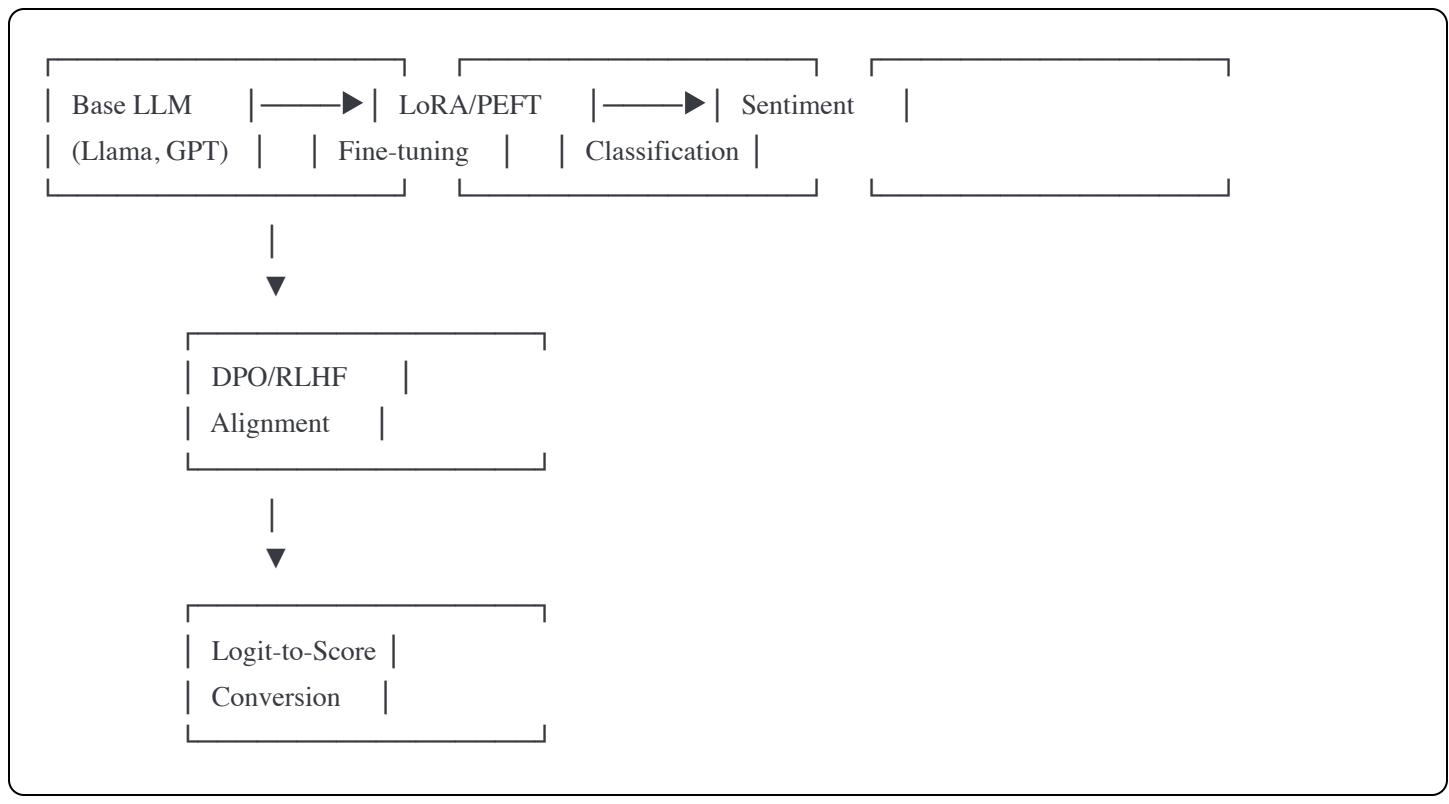
The FinRL Contest 2024-2025 demonstrated that ensemble methods effectively address policy instability by combining strengths of multiple RL agents—critical for volatile markets like crypto.

2.2 Large Quantitative Models (LQMs)

Meta Model Architecture (Numerai Style)

1. **Data Scientists**: Submit individual ML models predicting stock signals
2. **Staking**: Participants stake NMR tokens on predictions
3. **Aggregation**: Signals combined into unified Meta Model
4. **Execution**: Meta Model trades global equities
5. **Feedback**: Token burn/earn based on prediction accuracy

Financial LLM Fine-Tuning Pipeline



2.3 Zero-Knowledge Machine Learning (zkML)

Current State

zkML enables cryptographic verification of ML model outputs without revealing inputs, weights, or proprietary details.

Key Use Cases for Trading:

- Verify trading bot legitimacy without exposing strategy
- Prove model execution integrity on-chain
- DeFi pool rebalancing verification

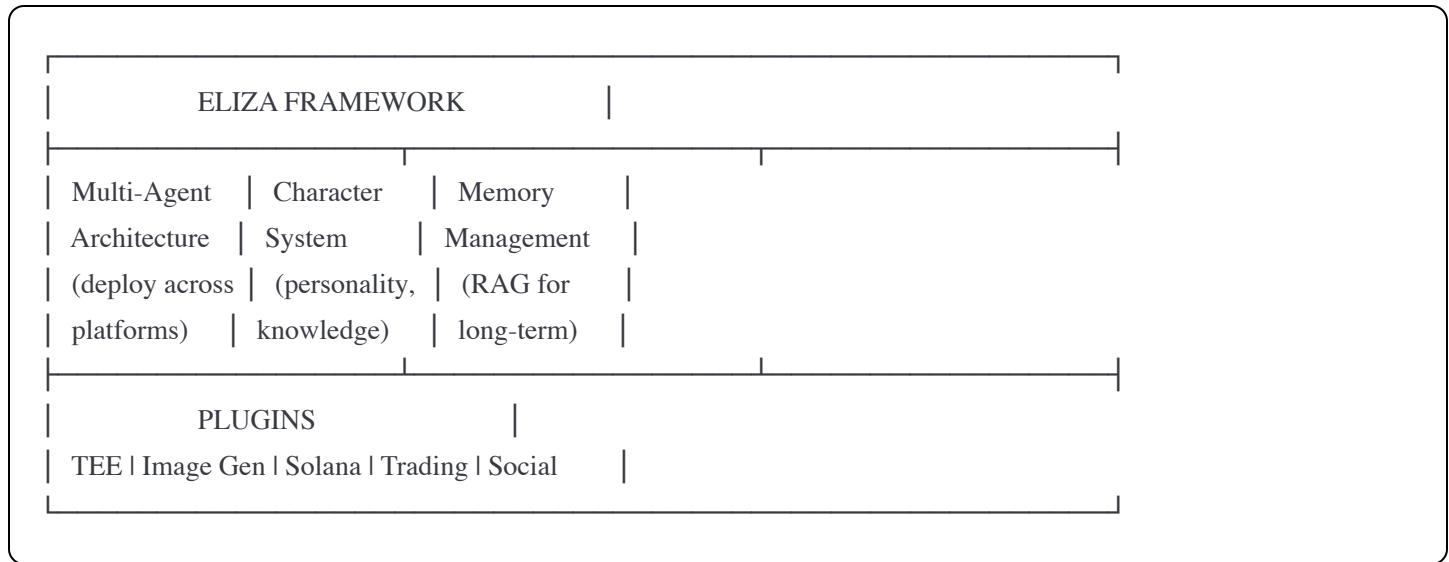
Projects Working on zkML:

- **Giza Protocol**: Original zkML focus, pivoted to semantic AI agents
- **Modulus Labs**: General zkML proofs
- **Worldcoin**: Identity verification via zkML
- **RISC Zero**: zkVM for verifiable computation

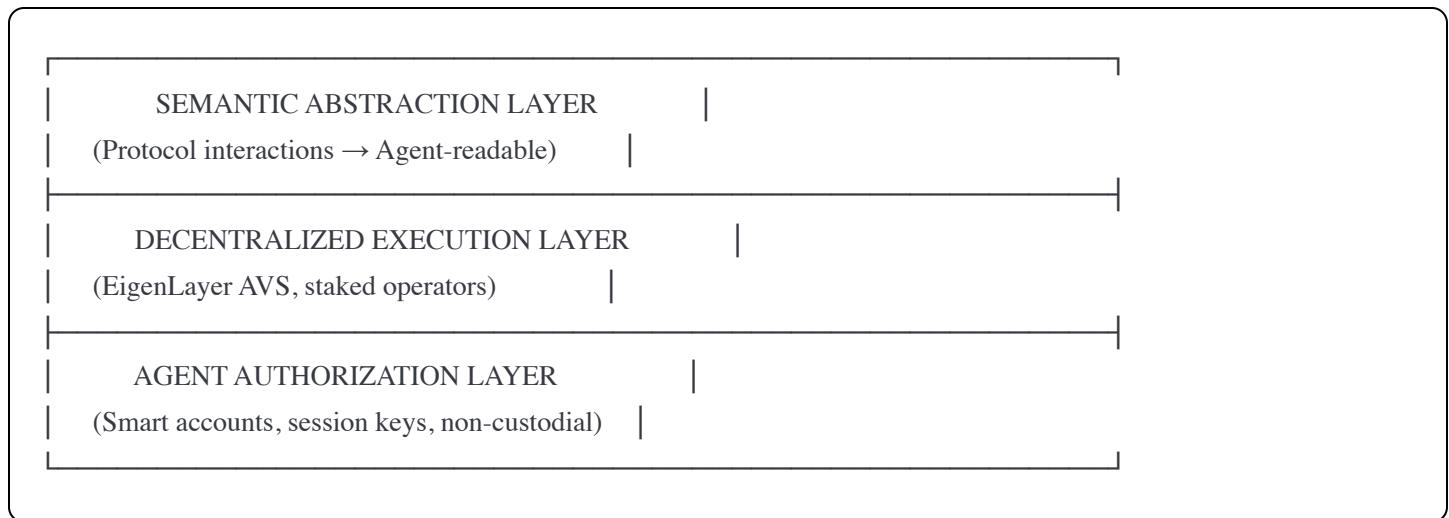
Limitation: Proving large models (GPT-4 scale) remains prohibitively expensive. Current implementations work for smaller trading bots and specific ML operations.

2.4 Agent Frameworks & Architectures

ElizaOS Multi-Agent Framework



Giza Protocol Stack



3. THE TOOLING & STACK (The "What")

3.1 Python Backtesting Libraries

Library	Speed	Best For	Live Trading
VectorBT	⚡ ⚡ ⚡ Fastest	Large-scale parameter sweeps	Via StrateQueue
Backtrader	⚡ ⚡ Medium	Full-stack development	Native IB/Alpaca
Zipline-Reloaded	⚡ Slower	Factor-based equity research	Via adapters
NautilusTrader	⚡ ⚡ ⚡ Fast	Institutional production	Native support
QuantConnect LEAN	⚡ ⚡ Medium	Multi-language (C#/Python)	Cloud deployment

Recommendation Matrix

- **Research-only strategies:** VectorBT or Zipline-Reloaded
 - **Retail trading systems:** Backtrader
 - **Institutional production:** NautilusTrader
-

3.2 Financial RL Frameworks

FinRL (AI4Finance Foundation)

- **GitHub:** github.com/AI4Finance-Foundation/FinRL
- **Features:** 30+ market data sources, vectorized environments, ensemble methods
- **Algorithms:** DQN, DDPG, PPO, SAC, A2C, TD3
- **Recent:** FinRL-DeepSeek integration for LLM-infused risk-sensitive trading

FinGPT

- **GitHub:** github.com/AI4Finance-Foundation/FinGPT
 - **Features:** Auto data curation pipeline, LoRA fine-tuning, sentiment analysis
 - **Models:** FinGPT v3.1-3.3 (ChatGLM, Llama variants)
 - **Cost:** ~\$300 per fine-tuning cycle (vs \$3M for BloombergGPT training)
-

3.3 Data & API Stack

Market Data Sources

Source	Data Type	Frequency	Cost
Yahoo Finance (yfinance)	OHLCV, fundamentals	Daily/Minute	Free

Source	Data Type	Frequency	Cost
Alpha Vantage	Technical indicators, FX	Real-time	Freemium
Polygon.io	Options, crypto, stocks	Tick	Paid
Alpaca	US equities	Real-time	Free tier
CCXT	100+ crypto exchanges	Real-time	Free

LLM/AI APIs

Provider	Model	Finance Use Case
OpenAI	GPT-4, o1	Sentiment, analysis
Anthropic	Claude	Structured financial reasoning
DeepSeek	Open-source	Cost-effective fine-tuning
Hugging Face	FinGPT variants	Self-hosted inference

3.4 Technical Analysis Libraries

```

python

# Core Stack
import pandas as pd
import numpy as np
import ta # pandas-ta for indicators
import talib # TA-Lib for 150+ indicators
from sklearn.ensemble import RandomForestClassifier
import xgboost as xgb
import tensorflow as tf
import torch

# Visualization
import plotly.graph_objects as go
import matplotlib.pyplot as plt

# Backtesting
import vectorbt as vbt
from backtrader import Cerebro

```

3.5 DeFi/Web3 Infrastructure

Agent Execution

Tool	Function
Coinbase x402	HTTP payment protocol for AI agents
Safe (Gnosis)	Multi-sig smart accounts
EigenLayer	Restaking for decentralized compute
Base	Low-cost L2 for agent transactions

Oracle & Data

Protocol	Use
UMA	Optimistic oracle (Polymarket)
Chainlink	Price feeds
Pyth	High-frequency price data

4. GAP ANALYSIS (The Opportunity)

4.1 Current Failure Modes

Latency & Execution

- **Problem:** AI inference time (100ms-2s) incompatible with HFT requirements (microseconds)
- **Impact:** Limits AI to medium-frequency strategies
- **Opportunity:** Edge-deployed, quantized models for sub-10ms inference

Hallucination in Financial Advice

- **Problem:** LLMs fabricate data, citations, and financial facts with high confidence
- **Stats:** AI hallucinations cost global economy ~\$67B in 2024; rates dropped 96% since 2021 but remain problematic
- **Impact:** Unreliable for autonomous execution without verification layers
- **Opportunity:** RAG + real-time data grounding + human-in-loop hybrid systems

Privacy & Strategy Leakage

- **Problem:** Centralized AI services expose trading strategies through API calls
- **Impact:** Alpha decay as strategies become public
- **Opportunity:** zkML for verifiable private inference; TEE-based execution

Overfitting & Regime Change

- **Problem:** RL agents overfit to training data; fail during market regime shifts
- **Impact:** Strategies that backtest well fail in production
- **Opportunity:** Ensemble methods, continuous retraining, regime detection

Regulatory Uncertainty

- **Problem:** AI trading faces unclear regulatory status (SEC, CFTC oversight gaps)
 - **Impact:** Institutional adoption hesitancy
 - **Opportunity:** Compliant infrastructure with explainability built-in
-

4.2 Market White Spaces

1. No "Institutional-Grade" Retail AI Trading

Current tools are either: (a) institutional-only (Renaissance), or (b) retail-quality with limited AI sophistication. There's no "Robinhood meets Numerai" product.

2. Cross-Asset Unified Agent

Existing agents are siloed: crypto-only or equities-only. No single agent intelligently allocates across stocks, bonds, crypto, and DeFi yields.

3. Explainable AI for Trading

Black-box models dominate. Regulators and risk managers need interpretable signals. Opportunity for "XAI Trading" layer.

4. Real-Time Sentiment at Scale

AIXBT and similar tools lag behind market moves. Opportunity for sub-second sentiment aggregation with verifiable accuracy.

4.3 Three "Killer App" Angles

ANGLE 1: The "Open-Source Numerai" for Retail

Concept: Decentralized hedge fund where retail users stake tokens on their ML models, aggregate into a meta-strategy, and share proportional returns.

Why Now:

- Numerai validated the model (\$550M AUM, 25% returns)
- ElizaOS/ai16z demonstrated retail DAO appetite (\$2B+ market cap)
- FinRL + FinGPT provide open-source ML infrastructure

Differentiation:

- Fully permissionless (Numerai is invite-only for staking)
- Multi-asset (stocks + crypto + DeFi yields)
- Mobile-first UX with guided model submission

Technical Stack:

- FinRL for strategy framework
 - Eigenlayer for decentralized execution
 - Giza-style semantic layer for DeFi integration
 - Token-gated governance via snapshot
-

ANGLE 2: The "Verifiable AI Trading Terminal"

Concept: Professional trading terminal where every AI-generated signal includes cryptographic proof of model integrity and data provenance.

Why Now:

- Institutional hesitancy around AI "black boxes"
- Regulatory pressure for explainability (EU AI Act, SEC scrutiny)
- zkML maturing for smaller model proofs

Differentiation:

- Every signal has zkML proof-of-inference
- Audit trail for compliance
- Hybrid human-AI decision support (not fully autonomous)

Technical Stack:

- RISC Zero / Modulus for zkML proofs
 - Fine-tuned FinGPT for sentiment signals
 - Traditional quant infrastructure (Bloomberg Terminal APIs)
 - Enterprise SSO and compliance dashboards
-

ANGLE 3: The "AI Agent Yield Optimizer" (Productized ARMA)

Concept: Consumer-facing app that deploys non-custodial AI agents to maximize yield across DeFi protocols—

no crypto expertise required.

Why Now:

- Giza's ARMA proved the concept (\$20M AUA, 15% USDC yields)
- DeFi complexity creates massive UX barrier
- "Set and forget" appeals to passive investors

Differentiation:

- Mobile-first, one-click deployment
- Insurance integration (Nexus Mutual, etc.)
- Fiat on-ramp directly to agent
- Natural language strategy customization

Technical Stack:

- Giza Protocol semantic layer
- Safe smart accounts for authorization
- Account abstraction (ERC-4337) for gas UX
- RAG-powered chatbot for strategy explanation

5. READING LIST

Academic Papers

Title	Focus	Link
FinLlama: LLM-Based Financial Sentiment Analysis for Algorithmic Trading	LLM + trading	dl.acm.org/doi/10.1145/3677052.3698696
FinRL Contests 2023-2025: Data-Driven Financial RL Agents	RL benchmarks	ietresearch.onlinelibrary.wiley.com
Enhanced Financial Sentiment Analysis Using LLMs	Sentiment trading	aclanthology.org/2024.wassa-1.1
Deep Reinforcement Learning for Automated Stock Trading	Ensemble strategies	arxiv.org/abs/2304.06037

GitHub Repositories

Repo	Description
github.com/AI4Finance-Foundation/FinRL	Financial RL framework
github.com/AI4Finance-Foundation/FinGPT	Open-source financial LLMs
github.com/Polymarket/agents	Autonomous prediction market trading
github.com/elizaOS/eliza	Multi-agent AI framework
github.com/polakowo/vectorbt	Fast backtesting library

Key Documentation

Resource	Link
Numerai Tournament Docs	docs.numer.ai
Giza Protocol Docs	docs.gizatech.xyz
FinRL Tutorial	finrl.readthedocs.io
FinGPT HuggingFace	huggingface.co/FinGPT
Virtuels Protocol Docs	virtuels.io/docs

Industry Reports

Report	Publisher
Big Ideas 2024: AI + Blockchain Intersection	Ark Invest
State of Crypto 2025	a16z crypto
AI in Asset Management	CFA Institute
Autonomous AI and Autonomous Agents Market Size	Global Market Insights

Newsletters & Analysis

Source	Focus
Bankless (bankless.com)	DeFi + AI agents
The Defiant	DeFi news
Analyzing Alpha (analyzingalpha.com)	Python trading tools
AI4Finance Blog	Academic FinML

6. KEY METRICS TO WATCH

Metric	Current (Nov 2025)	Signal
AI Agent Token Market Cap	\$4.34B	Sector health

Metric	Current (Nov 2025)	Signal
Numerai AUM	\$550M	Institutional validation
Polymarket Monthly Volume	\$1.5B	Prediction market growth
Giza AUA	\$20M	DeFi AI adoption
FinRL GitHub Stars	10K+	Developer interest
Virtuals Agents Launched	21,000+	Creator ecosystem

CONCLUSION

The AI x Trading landscape has moved from experimental to operational. Numerai proved crowdsourced AI can deliver institutional returns. Giza demonstrated autonomous DeFi yield optimization is viable. ElizaOS and Virtuals showed retail demand for AI agent ownership.

The opportunity for a next-generation trading solution lies in the gaps: combining institutional-grade AI with retail accessibility, adding verifiability to black-box models, and unifying fragmented cross-asset strategies into coherent agent systems.

The "picks and shovels" are available: FinRL for reinforcement learning, FinGPT for sentiment analysis, VectorBT for backtesting, and Giza/Eigenlayer for decentralized execution. The challenge is integration, UX, and regulatory navigation.

Recommended Next Steps:

1. Prototype with FinRL + VectorBT on historical crypto data
 2. Fine-tune FinGPT v3.3 on your target market's news corpus
 3. Deploy test agent on Base using Giza's semantic layer
 4. Validate signals against Numerai tournament data
 5. Build compliance framework before scaling
-

Dossier compiled November 2025

For strategic planning purposes only – not financial advice